

# A Scalable, Flexible and Configurable OOM Killer

Nikhil Abhyankar  
*nikhilsa@vt.edu*

Gaurang Karwande  
*gaurangajitk@vt.edu*

Radhika Nibandhe  
*radhikan@vt.edu*

## 1 Introduction

Every program running in a computer system accesses memory in order to execute and generate results. The Linux Memory management subsystem is responsible for managing the memory. The work of this subsystem is not restricted just to allocating memory for user programs and kernel internals but also implementing virtual memory and demand paging. When a system is utilizing virtual memory to allocate memory to a user space program, the correspondence between virtual and physical memory is setup only after the first attempt to access the memory is made. Overcommit refers to the practice of giving out virtual memory with no guarantee for the committed amount of physical storage to exist. This is done to avoid the under utilization of the RAM. The overcommit can thus be harmful and cause the system to crash. The kernel is designed to perform memory overcommit handling heuristically by estimating the total available memory thus terminating requests for large chunks of memory. As a result of such a heuristic process, memory overload is possible. The overcommit policy in the kernel is set via the `sysctl vm.overcommit_memory`.

Out of Memory (OOM) is an error message which occurs when additional memory allocation is not possible for programs or the operating system. Thus, Out of Memory is an undesired state in a computer system. An OOM can error causes programs or even the entire computer system to shut down. The computer forcibly shuts down programs or the computer.

The kernels of operating systems try to recover from such an OOM condition by terminating processes. If an application dereferences the pointer and if the system has no memory space available, then the OOM Killer is invoked [1]. Thus, the OOM Killer is responsible to prevent the running processes from unitedly wearing out the host memory. Whenever a process tries to allocate more than the available memory, the process having the highest badness score will receive an out-of-memory signal,

which dictates the process to quit or terminate one of its child processes. If neither is done, the process is terminated spontaneously by the OOM Killer. However, the controversial part is that, the process which is terminated may not be the same as the process triggering the OOM signal. It chooses a process that is using large memory but not long lived as a general assumption is made that long lived processes are not affecting the memory allocation process.

The victim selection mechanism of the kernel's OOM killer has been a point of contention in recent times. As it may select a victim process using more memory relative to other processes but at the same time is more critical than other processes running in the background. A very good example of such a process is the training and inference phase of a Machine Learning model. Machine Learning models are becoming more and more resource heavy with the recent success of Neural Networks. Deep Learning networks need a lot of data along with complicated model architectures which have resulted in tremendous need for resources being allocated to them. Currently, GPUs are used to handle this, however, an emerging trend is containerizing the ML models to solve the problem. Containerizing of the ML workflow requires putting the ML models in a container and then deploying it on a machine. With the increase in the number of applications of ML, the users have also gone up. Increased users containerisation has led to users migrate towards common cloud clusters for executing their works[2][3]. This can lead to a problem of shortage of resources on the servers hosting ML workflows requiring heavy computation. In case there is a memory over commit, some of the processes will be killed, often bringing down containers with them.

## 2 Motivation

The kernel OOM killer works on self-defined heuristics to select the victim process. The function

`select_bad_process()` chooses the process to be terminated [4]. Put very simply, it selects the process with the highest `oom_score`. How this `oom_score` is calculated has been a highly contended topic and has received fair amount of criticism on various online forums. The kernel calculates the `oom_score` for each individual process, and it can be viewed by running the following command -

```
$ cat /proc/<process-PID>/oom_score
```

The `select_bad_process()` acts like a debugger going over each running task to calculate the `badness()` which is responsible for killing the process. The badness is calculated as follows, note that the square roots are integer approximations calculated with `int_sqrt()`.(figure1)

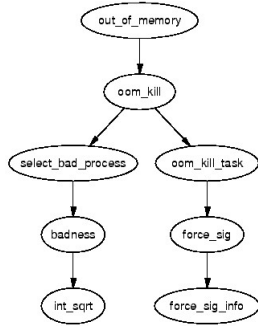


Figure 1: OOM Kill Graph [4]

The OOM Killer is invoked when the memory pressure reaches desperate levels. Its job is to kill one of the existing processes. Every process has an `OOM score_adj` value that is used to tweak the score relative to that of other processors. The OOM victim selection process is unpredictable as it may end up killing the processes in the user space than the kernel space to maintain system stability.

$$\begin{aligned}
 \text{badness\_for\_task} \\
 &= \text{total\_vm\_for\_task} / (\text{sqrt}(\text{cpu\_time\_in\_seconds}) \\
 &\quad * \text{sqrt}(\text{sqrt}(\text{cpu\_time\_in\_minutes}))) [4]
 \end{aligned}$$

### 3 Related Work

Linux users have been occasionally troubled by the misbehaving OOM Killer by killing any random process. Facebook (now Meta) has been one of the biggest critiques of the OOM Killer as it has become a hindrance to their humongous servers running primarily on Linux[1]. However, the OOM Killer cannot be configured at the kernel level and hence a user-space alternative is necessary[5].

As a result, Facebook came up with OOMD [1], which is

a OOM daemon for the User-Space with an idea of providing more flexibility to the kernel code. The daemon is triggered when the system is under pressure and can be invoked whenever certain threshold levels are breached. OOMD uses PSI and `cvgroupv2` for system monitoring. OOMD takes the required action before the kernel space OOM is active. A custom code can be written to make the system configurable enabling user defined protection rules for safeguarding the running processes.

The EarlyOOM implementation checks the amount of available memory instead of free memory and memory available for swapping. By default, if both the memory values drop below 10%, then the the OOM killer is activated and the process with the highest OOM score is selected to get killed by the OOM killer. A similar implementation along with additional features has been done in Python.[6]

The work proposed in this project is inspired from the above contemporary works and aims to incorporate user-defined policies in the OOM Killer's victim selection mechanism. While we are inspired by memory consumption trends for the training loop of an ML model, we design a system that is modular and scalable to any processes deemed to be critical by the user.

## 4 System Design

The project is divided into three parts -

1. **Kernel Module:** To monitor system memory statistics
2. **OOM Daemon:** User-space program that handles victim selection and killing
3. **Test Programs:** To test and evaluate the OOM daemon

### Kernel module

A Kernel module is an external segment of code that can be unloaded and loaded into the kernel as and when required without requiring to reboot the system. The first step in the project is to gather system memory statistics and present it to the user. A kernel module is thus created to monitor memory statistics information such as Total System Memory, the proportion of Available Memory, Total Swap Memory, proportion of Available Swap Memory and the per process individual memory. This information forms the basis of activating the User-space OOM Killer. The existing implementation raises a flag to summon the OOM daemon.

```
1. → ~ ls -l /proc/pressure
2. total 0
3. -r--r--r-- 1 root root 0 Aug 28 07:06 cpu
4. -r--r--r-- 1 root root 0 Aug 28 07:06 io
5. -r--r--r-- 1 root root 0 Aug 28 22:22 memory
6. → ~
```

Pressure metrics along with other tools help detect the resource shortages, killing the non essential processes and reallocating memory to the system. PSI provides fair warnings of resource shortages and warns the user beforehand, in case the memory starts becoming scarce. [9]

some: 50.00  
full: 16.66

60 seconds

The fields `cpu`, `memory`, `io` give corresponding memory information. The output of `/proc/pressure/cpu` gives the average times processes were starved of CPU. The output for PSI memory is shown in figure 4.

```
→ ~ cat /proc/pressure/cpu
some avg10=0.05 avg60=0.04 avg300=0.01 total=576927401
```

Percentage of the time on average every 10, 60 and 300 seconds processes were starved of CPU.

Total time in microseconds processes were starved of CPU

```
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=10.130000
full_avg10=8.300000
full_avg10=5.560000
full_avg10=3.730000

[i] 0:~bash*
```

## OOM Daemon

The work presented focuses on achieving the communication between the Kernel module and user-space OOM killer. The victim selection module of the user-space OOM killer solves an optimization problem with the objective of releasing the most memory pressure and at the same time achieving the best overall system utility. Thus we formulate the victim selection module as a variant of the Knapsack Problem. Dynamic programming techniques such as knapsack have been widely used in process scheduling in the CPU.[14]

3

rived from the idea of filling a sack of fixed size with the most valuable items [12]. In our case, the weight of a process is its OOM score and the value is a user-defined quantity which we call the process utility. The maximum weight threshold is also a user-defined quantity, however this can easily made dynamic by incorporating some heuristics about trends followed by individual process OOM scores.

## Test Programs

In this module we run a batch of test processes to evaluate our implementation. Each process has an user assigned process utility. In simple terms, process utility is how much benefit or value we are getting by allowing that process to run through completion. We have created a replica of the `/proc` directory as `/tmp/user_processes` where information about the test processes are stored. The OOM daemon can access all the necessary information from this location.

## 5 Experimentation

The first phase of our experimentation work consisted of creating a OOM daemon to kill a process based on user-defined priority. Figure.10 shows the working of our user-space OOM killer when tested with two processes. Here we run two test processes 4904 and 4905 with priorities 100 and 20 and OOM score of 983 and 1150. Smaller the priority more critical is the process. These priorities are assigned by the user. We can then see that even though process 4904 has lesser OOM score, it is killed first as it is deemed to be less critical by the user. Then process 4905 is killed once the OOM condition is recreated.

As our initial motivation has been about taming the kernel's OOM killer for high memory utilization workloads like deep learning, we run a batch of processes consisting of computer vision tasks and some dummy processes. Computer Vision tasks involve complex matrix multiplications and are computationally quite heavy. The idea is to ensure smooth functioning of the ML processes with the heavy memory consuming dummy processes getting killed. We run a DeepFake image generation process, based on an AutoEncoder, called FaceSwap wherein the target face is superimposed on the source face over multiple epochs to create a fake image. The network consists of two decoders, one each for target and source and one common encoder. The network learns the common features of the source and target from the encoder by storing it in a latent variable. This is then extrapolated using the decoder to generate the face back. [13].

More precisely, we execute two FaceSwap processes which are Machine Learning Processes with more than




Figure 6: FaceSwap process output

one hundred thousand epochs each and two dummy processes. These dummy processes are designed to continuously allocate memory, hence available system memory and PSI metric quickly reach their thresholds and user-space OOM killer is activated. Dummy process with PID 1609 has utility of value 8 and OOM Score 916. Dummy process with PID 1610 has utility of 5 and OOM score 778. The ML process with PID 1607 has utility between 50 to 100 which is randomly generated and OOM Score 746. ML process with PID 1608 has a utility value that lies between 50 to 100 and is randomly generated and OOM Score 673. Figure.7 shows the test processes running in tandem. The first quarter is the dummy process 1609. Second quarter is the first ML process 1607. The third quarter is the dummy process 1610. And the last quarter is the second ML process 1608.

```
[radhikan@fedora 09_New_OOMKILL]$ cd test_proc1
[radhikan@fedora test_proc1]$ ./userOOMtest
Test process PID: 1609
|| 19 MiB || 2874 MiB/s ||
|| 39 MiB || 2957 MiB/s ||
|| 58 MiB || 2955 MiB/s ||
|| 78 MiB || 2795 MiB/s ||
|| 97 MiB || 3025 MiB/s ||
|| 117 MiB || 2882 MiB/s ||
|| 136 MiB || 2081 MiB/s ||

[radhikan@fedora 09_New_OOMKILL]$ cd ML_Proc1
[radhikan@fedora ML_Proc1]$ python train.py
====> Using CPU to train
1607
====> Loading datasets
Killed
[radhikan@fedora ML_Proc1]$

[radhikan@fedora 09_New_OOMKILL]$ cd test_proc2
[radhikan@fedora test_proc2]$ ./userOOMtest
Test process PID: 1610
|| 19 MiB || 1684 MiB/s ||
|| 39 MiB || 3055 MiB/s ||
|| 58 MiB || 1188 MiB/s ||
|| 78 MiB || 914 MiB/s ||
|| 97 MiB || 844 MiB/s ||
|| 117 MiB || 916 MiB/s ||
|| 136 MiB || 66 MiB/s ||

[radhikan@fedora 09_New_OOMKILL]$ cd ML_Proc2
[radhikan@fedora ML_Proc2]$ python train.py
====> Using CPU to train
1608
====> Loading datasets
====> Start from scratch
Start training, press 'q' to stop
epoch: 0, loss:0.18196828663349152, lossB:0.18559400737285614
"fedora" 15.42 03-Dec-21
```

Figure 7: Test Processes

Figure.8 shows the output of the created kernel module. This module continuously monitors the system and outputs the Available Memory and Swap Memory usage and accordingly sends the signal to activate the user-space OOM Killer depending on comparison made with

the threshold values. The current threshold is set at 30% of the total available memory.

PID	Memory Available in System	Memory Available for Swap
+0.502685	90%	100%
+0.503605	90%	100%
+0.505256	90%	100%
+0.507139	90%	100%
+0.513098	90%	100%
+0.513310	90%	100%
+0.509060	90%	100%
+0.511104	90%	100%
+0.508807	56%	100%
+0.504140	4%	100%
+0.503843	0%	89%
+0.503987	0%	76%
+0.503984	0%	62%
+0.503998	0%	49%
+0.503993	0%	35%
+0.505013	0%	21%
+0.503189	23%	81%
+0.503785	87%	98%
+0.517857	80%	98%
+0.506149	63%	98%

Figure 8: Memory Statistics: Kernel Module

```

[radhikan@fedora PSI_Monitor]$ ./PSI_Mem
poll_interval = 5s,
recovery_interval = 15s,
stall_threshold = 30%

full_avg10=0.000000
full_avg10=3.980000
full_avg10=33.580002

```

Figure 9: User Space:PSI Implementation

In addition to system memory levels we also incorporate the Pressure Stall Information (PSI) while deciding when to activate the user-space OOM killer. The implementation of the PSI in the system can be seen in figure. 9. The threshold is set at 30% to summon the OOM daemon earlier. The recovery interval the time for which the system forces all other processes to sleep while the existing process is getting killed. The Polling will take place to read the memory from proc file to check the pressure on memory.

## 6 Results

Figure.11 showcases the output of the User Space daemon along with PSI executed using a bash file `testproc.sh`. The execution includes running of 2 ML processes along with 2 dummy processes. However, as a result of running the processes using a bash script, the output is consolidated, not including the view for all processes separately. The victim selection module within the OOM daemon selects a batch of processes to kill such that it results in the overall highest user assigned utility and releasing the most memory pressure. Initially, we send a SIGTERM signal to the victim process. This gives the victim process to correct or reduce its memory use. We have purposely blocked the SIGTERM signal in all our test processes. Next the system sends SIGKILL to the victim process and the process is instantly terminated. We see that the dummy processes are killed first. Their utilities are 5 and 8, and their OOM scores are 674 and 1132. respectively. They have low utility (value) and a bad OOM score (weight). We set the OOM score threshold to 1000. Hence the victim selection module tries to select victim processes such that the summation of all running processes' OOM score is kept below this threshold. We also see that one of the non-dummy process is killed. Process 1115 is one of the two ML processes. It's utility is between 50 to 100 since the utility is randomly generated. And the OOM score at the time of killing is 906. Hence among the processes of interest it is a relatively 'bad' process. We see that after the first run of OOM Killer, the overall system OOM Score is reduced by 906, at the expense of a utility cost of 5. Then when the OOM condition is recreated the OOM Killer is invoked again. This time it results in the reduction of the OOM Score by 1806 at the expense of a utility cost of 98.

In Figure.8, we can see that the available System Memory and Swap Memory go on decreasing. The System Memory reaches 0% and the Swap Memory reaches 21%. Hence the memory falls below threshold and our OOM killer is invoked. After the execution of the OOM killer we can see that the memory levels increase and are back above the threshold.

Figure.11 displays a PSI Implementation for the pressure on memory. The display shows the PSI threshold which has been set to 30%. This threshold represents the value the processes have been on hold for 10 seconds. The display will print the full avg10 notifying the user and system about the pressure on the system every 10 seconds.

Hence, the user-space OOM killer is successful in selecting and killing processes based upon a user-define custom policy. We are effectively doing Out Of Memory management in user-space and the kernel's OOM

killer is stalled from coming into play. We have in no way changed the kernel's OOM implementation and the kernel's OOM killer will still work if our system fails to successfully release the memory pressure. However based upon our design this scenario will never arise.

## 7 Project Goals

The initial project goals were -

1. **75% Goal:** Create a module to monitor the current memory usage statistics and send activation signal to user-space OOM killer.  
**Status:** Completed.
2. **100% Goal:** Victim process selection by user-defined custom policy  
**Status:** Completed.
3. **125% Goal:** Formulation of the Knapsack problem and testing the User Space OOM killer with a batch of ML workloads and other non-critical processes  
**Status:** Updated the goal after experimentation.

We initially envisioned that the kernel module would be able to classify if a process is an ML process based upon the memory allocations trends. But after experimentation we realized that this was not the case. We tried monitoring the memory trends - the memory pressure, first and second gradients of the memory allocation, for making the classification. But the logic was too complex and not accurate. Hence the current OOM Killer that we have designed works irrespective of the type of the process. It can be extended to any process the user deems important and doesn't want to be killed by the Linux implementation of Out of Memory management.

## 8 Tasks Accomplished

1. Study of memory management techniques of the Linux kernel along with the implementation of the OOM Killer in the kernel-space.
2. Creation of a kernel module in order to periodically monitor the system memory statistics to result in generation and delivery of an activation signal to the User space OOM killer, based on comparison with a set threshold.
3. Implementation of the OOM daemon in the User Space to emulate the OOM Killer in the Kernel space. The OOM daemon is invoked before the Kernel space OOM Killer making the process flexible and configurable.

4. Customising the way to kill a process in the OOM daemon based on user-defined process utilities.
5. Implementation of PSI along with memory statistics information to summon the User space OOM daemon.
6. Formulation and implementation of an Knapsack Optimisation algorithm to decide the process to be killed.
7. Creation of an ML based resource (memory) heavy process and dummy test processes for checking the successful execution of the OOM killer.
8. Creation of a bash script to enable running all the processes simultaneously.

## 9 Conclusion and Future Work

The proposed system monitors the memory allocation and PSI for the processes and selects a process based on Knapsack Optimisation, such that the utility as well as the memory pressure released is maximized. The system successfully customises the killing of a process in the user space as against the kernel's controversial 'arbitrary' OOM killing.

The heuristics for victim selection is based upon the conventional knapsack problem. However this requires the process utilities to be defined by the user. This utilities can be instead computed dynamically based upon various factors such as the rate of change of a process OOM score and its individual memory utilization trends. We plan to investigate this further. Our initial motivation was to design a kernel module which can distinguish between an ML and non-ML workload. This however requires a large enough labelled training data-set. The experiments conducted suggest that availability of large dataset related to memory statistics, memory pressure and memory allocation can facilitate the use of Machine Learning to select a process to be killed. Additionally with the recent advent of Reinforcement Learning, a reward system can be designed in such a way that the OOM killer does victim selection based upon the past user trends. This can circumvent the need for a large labelled dataset. We plan to investigate this more in future work.

## 10 Member Contributions

- Nikhil:  
Making Memory Tracker Module, Making User Space OOM Killer, Adding Custom Policies, Literature review, Testing with ML processes, Report Writing

- Gaurang:  
Formulation and understanding of the problem, Understanding Kernel's OOM Score, Making User Space OOM Killer, Code Design, Report Writing, Literature Review, Debugging, Knapsack Optimisation for Process Killing
- Radhika:  
PSI information utilization and implementation, Making User Space OOM Killer, Adding Custom Policies, Debugging, Report Writing, Literature Review

- [14] Neelakantagouda Patil. "A Knapsack Based CPU Process Scheduling Using Neelsack Algorithm". In: (2015), pp. 138–145. URL: <http://ijseas.com/volume1/v1i7/ijseas20150715.pdf>.

## Acknowledgement

The authors of the report would like to thank Dr. Changwoo Min and Mr. Christopher Jelesnianski for their guidance and support throughout the Fall 2021 semester for the 5414G: Advanced Linux Kernel Programming course.

## References

- [1] URL: <https://github.com/facebookincubator/oomd>.
- [2] URL: <https://www.bmc.com/blogs/machine-learning-containers/>.
- [3] URL: <https://www.instana.com/blog/solving-the-out-of-memory-killer-puzzle/>.
- [4] URL: <https://www.kernel.org/doc/gorman/html/understand/understand016.html>.
- [5] URL: <https://github.com/rfjakob/earlyoom>.
- [6] URL: <https://github.com/hakavlad/nohang>.
- [7] URL: <https://www.kernel.org/doc/html/latest/accounting/psi.html>.
- [8] URL: <https://unixism.net/2019/08/linux-pressure-stall-information-psi-by-example/>.
- [9] URL: <https://facebookmicrosites.github.io/psi/docs/overview>.
- [10] URL: <https://lwn.net/Articles/590960/>.
- [11] URL: <https://lwn.net/Articles/552789/>.
- [12] URL: [https://en.wikipedia.org/wiki/Knapsack\\_problem](https://en.wikipedia.org/wiki/Knapsack_problem).
- [13] URL: <https://github.com/Oldpan/Faceswap-Deepfake-Pytorch>.



## Appendix

```
[radhikan@fedora UserSpaceOOM]$ ./userOOM
Process 4905 priority parsed
Process 4904 priority parsed
Process 4811 priority parsed
/proc read error
Sending SIGTERM to process: 4904
OOM_score 983
VmRSS 2286772
Priority: 100
Process 4904 priority file deleted
Process 4904 killed

Process 4905 priority parsed
Process 4811 priority parsed
/proc read error
Sending SIGTERM to process: 4905
OOM_score 1150
VmRSS 4522728
Priority: 20
Process 4905 priority file deleted
Process 4905 killed

[radhikan@fedora test_proc1]$
[radhikan@fedora test_proc2]$

Killed
[radhikan@fedora test_proc1]$
[radhikan@fedora test_proc2]$

Killed
[radhikan@fedora test_proc2]$
```

Figure 10: Output of running two dummy processes

```
radhikan@fedora:~/share/01_Share/09_New_OOMKILL
Process 1114 priority parsed
/proc read error
Sending SIGTERM to Process: 1115 OOM_score 906 Value: 5 VmRSS: 2186276
Process 1115 priority file deleted
Process 1115 killed

Freed totals:          906      5
Process 1113 priority parsed
Process 1112 priority parsed
Process 1114 priority parsed
/proc read error
Sending SIGTERM to Process: 1113 OOM_score 674 Value: 90 VmRSS: 95408
Process 1113 killed

Sending SIGTERM to Process: 1114 OOM_score 1132 Value: 8 VmRSS: 4361940
Process 1114 priority file deleted
Process 1114 killed

Freed totals:          1806      98

Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=12.080000
Full_avg10=9.890000
Full_avg10=6.630000
Full_avg10=4.440000
Full_avg10=2.440000
Full_avg10=1.340000

[0/3] 6953 MiB | 937 MiB/s |
6972 MiB | 757 MiB/s |
6992 MiB | 828 MiB/s |
7011 MiB | 988 MiB/s |
7031 MiB | 908 MiB/s |
7059 MiB | 821 MiB/s |
7078 MiB | 1090 MiB/s |
7089 MiB | 798 MiB/s |
7109 MiB | 812 MiB/s |
7128 MiB | 811 MiB/s |
7148 MiB | 868 MiB/s |
7167 MiB | 903 MiB/s |
7187 MiB | 1026 MiB/s |
7207 MiB | 997 MiB/s |
7226 MiB | 795 MiB/s |
7246 MiB | 722 MiB/s |
7265 MiB | 580 MiB/s |
7285 MiB | 816 MiB/s |
7304 MiB | 1066 MiB/s |
7324 MiB | 1544 MiB/s |
7343 MiB | 856 MiB/s |
7363 MiB | 847 MiB/s |
7382 MiB | 556 MiB/s |
7402 MiB | 1420 MiB/s |
7421 MiB | 875 MiB/s |
7441 MiB | 958 MiB/s |
7460 MiB | 814 MiB/s |
7480 MiB | 869 MiB/s |
7500 MiB | 915 MiB/s |
7519 MiB | 837 MiB/s |
7539 MiB | 780 MiB/s |
7558 MiB | 789 MiB/s |
7578 MiB | 921 MiB/s |
7597 MiB | 757 MiB/s |
==> Start from scratch
Start training, press 'q' to stop
epoch: 0, lossA:0.18734067678451538, lossB:0.1763179451227188
==> Saving models...
epoch: 1, lossA:0.19053064286708832, lossB:0.17658859491348267

[0] 0:[tmux]
```

Figure 11: Output of running ML processes along with a dummy processes